

memory

Chhi'med Kunzang

July 27, 2019

Contents

1 ZigZag Replication Memory Requirements (as multiple of sector size)	1
1.1 Why does it matter?	1
1.2 Merkle trees	1
1.3 Encoding	2
1.3.1 Idea: Ensure parents need not be accessed randomly. .	2

1 ZigZag Replication Memory Requirements (as multiple of sector size)

1.1 Why does it matter?

- RAM is expensive and requires expensive configurations at high levels.
- We consider 64GiB the limit for ordinary operation.
 - Someone might always prefer more, but we ideally will not require it.
- Larger sectors are more economical from both a CPU and proof size perspective.

1.2 Merkle trees

- Most naively, we need enough memory to hold all trees at once.
- Improve by offloading trees to disk and reload for proofs.
- Reduce space usage per tree.
- Best memory usage holds \log_2 nodes but uses more CPU.

1.3 Encoding

- We want replication time to be strictly limited by linear encoding time.
- This means no waiting for parents to load.
- Parent distribution is random and resists locality.
- Hashing (currently blake2s) a parent for the KDF is faster than random-access latency on disk.
- This seems to imply we need to keep the entire replica in memory, if we don't want to sacrifice speed (which we don't).
- If we could take advantage of sequential reads, using disk might be fast enough to outrun hashing.

1.3.1 Idea: Ensure parents need not be accessed randomly.

- This implies greatly increased total space, since each parent has (on average) P (currently 13) children.
- Two approaches using multiple machines:
 - Allow replicating sectors larger than 64GiB but still requiring at least as much RAM as the sector size (on multiple machines).
 - Allow replicating sectors larger than the total RAM.
- Start with the first, less ambitious target and build up to the least bounded we can.

1. The way that doesn't work

- Just write each parent to the ~ 13 locations optimal for encoding each of its children.
 - Unfortunately, this trades random reads for random writes, which are if anything more expensive.

2. Divide responsibility for one sector across multiple machines.

- If writing to disk, we **can** do so in parallel. With enough machines, this makes the amortized cost of writing fast enough.
- If holding in memory we can replicate sector sizes greater than our (64 GiB) RAM limit on any given machine.

- We will need much more total memory across all machines (order of $\times P$) than the sector size, though.

3. Algorithm

(a) Basic Operation

- Each unit is responsible for encoding N consecutive nodes and runs on a separate machine.
- Layout is $[[\text{Node1} \dots \text{Parents1}] \dots [\text{NodeN} \dots \text{ParentsN}]] = N \times ((1 + P) \times 32)$ bytes per unit.
- Units are arranged in a ring, such that each unit has a 'next unit'.
- Units are ordered such that consecutive units are responsible for consecutive ranges of nodes.
- After encoding a node, the encoding unit sends the node's index and new value to the next unit and also processes it.
 - The next unit forward the message to its next unit and begins processing.
 - This continues until all units have seen the message.
- Each unit searches its layout for locations where the new node occurs in a 'parent' location, and stores the value there.
 - Naively, this requires the parent indexes to be colocated with the data. We will eliminate this later.
 - The originating unit also searches and writes the new value anywhere necessary in its remaining (future) layout.
- At any time, one unit is responsible for encoding the current node.
 - Assuming all messages have been processed, it is guaranteed that all of the current parents have been received.
 - Therefore, the current parents exist in contiguous memory and can be looked up directly.
- When a given unit reaches the end of the range of nodes for which it has responsibility, encoding passes to the next node.
- Upon reaching the end of a layer, the direction of encoding reverses, and the process repeats with order of units and nodes reversed.
 - This requires a change of layout for each unit.
- NOTE: There is clearly room for optimization, since with all data held in RAM, we do not need it to be sequential.

- We will offload some of this contiguous memory to disk.
- (b) How are matching parent locations found?
- We can precompute a correctly-ordered map and store it alongside the node-parents layout.
 - To precompute: record the locations (in layout coordinates) of each parent in layout.
 - If parent is not contained in unit's layout, no entry for parent is included in the map.
 - Maintain a pointer/index to the current map entry, advancing as each matching node is processed.
 - As each node is processed, its corresponding map entry can be immediately checked. If no match, do nothing.
 - If the processed node matches the current map entry, store the node in each of the entry's parent locations, and advance the entry pointer.
- (c) How do we reduce memory?
- As described here, more than P (13) times the sector size is required to replicate one sector.
 - Since both the map and the layout eventually contain correctly ordered sequence data, they are well-suited to non-random reads.
 - We can reduce total memory by choosing a percentage of both the map and the layout to store on disk.
 - Writes to the memory-resident portion remain fast.
 - Writes to the currently disk-resident portion become random-access and are slow.
 - However, since these writes are spread across multiple machines, each with its own disk, the cost of writes can be amortized.
 - Define one random-access write to be a factor of W slower than one KDF hash (ignoring the extra hash for the replica id).
 - Define the fraction of each unit's map/layout which is stored to disk as D .
 - For each node processed, we make P writes.
 - Of these, $D * P$ are random-access.

- In order not to be slower than encoding, we need that $D * P * W \leq 1$.
 - For example, suppose random writes are 5 times slower than KDF hashes. Then $W = 1/5$, $P = 13$, and $D \leq 5/13$.
 - That is, we can cache up to 0.385 of each unit.
- (d) Is it worth it? And if so, at what scale?
- For now, estimate that we can cache 1/3 of each unit's data.
 - Concrete example:
 - Ignoring the cost of the map, for now, total memory for 1TiB would be 14TiB.
 - Caching 1/3, we are left with 9.33TiB, or roughly 145 * 64GiB machines.
 - Assume this allows us to uninterruptedly encode.
 - Then we are able to replicate 1TiB using the same amount of RAM as could otherwise be used to replicate 9.33TiB.
 - In order for this enormous overhead to be worthwhile, we would need to see a 9.33X benefit in combined proof size and proving CPU time savings.
 - Proof size scales linearly with sector size, proving time logarithmically.
 - Therefore, although we pay a premium of 9.3X, we save 145X in proof size, and 20X in proving CPU time.
 - [Since we ignored and did not estimate the cost of the map, these numbers may be off by a relatively small factor.]
 - So, at scale, maybe it is worth it. In this model, practical sector size is determined by replication speed and desired turnaround time for sealing.
- (e) Further considerations
- Estimate the size of the unit's map.
 - NOTE: even assuming calculations are correct, all numbers above are based on the hypothetical and unresearched 5x difference between one KDF hash and one random write.
 - Which parts of the map/layout are cached, and how?
 - Since we want current reads to be fast, cache the last, not first portion.
 - As replication proceeds, we can reclaim memory which will not be needed again for this layer.

- Periodically flush the encoded nodes' map/layout to disk and replace it with disk-resident map/layout.
- At this point, we're manually managing disk paging.
- Our data is no longer random access.
- Can we just mmap the whole (correctly sized) data and let the system's virtual memory manage the problem for us?
 - * Maybe not, since the random writes may wreak havoc with this. Needs better understanding.